

Midterm Exam Review

Note Title

Production Systems & Problem Solving

Formalizing a "problem"
and a solution approach

Problem : State Representation — Initial State

Knowledge Base

Rules ("Moves")

- precondition : is Rule r applicable to state s ?
- action : what is the result when Rule r
is applied to state s ?

Goal condition goal(s) : Does state s satisfy goal condition?

Note: For a given state, many rules may be applicable

A solution strategy must decide which rule to apply when many are applicable.

Solution Strategies : Irrevocable , Tentative

↓
Keep track of current state only; applying rule r to state s produces new value of state s . When algorithm completes, state s is the solution state.

↓
Remember what has been done up til now; rules that have been applied + didn't work may be revoked — try something new, etc

Implicit Graph : Begin with start Node, build graph by choosing a node and generating successors with a successor Function (e.g., "ApplyMove(rule, s)")

a. Flailing Wildly

b. Hill Climbing

- variations

- simulated annealing

irrevocable

use evaluation function $f(\text{state})$

by never choosing a state with a lower evaluation than its predecessor, we ensure we will not return to any previous state.

allow algorithms to choose "best neighbor" even if that value is lower than current state

..... potential problems

- try from many different starting positions
- use different "tie-breaking rules", etc

However, we may get "stuck" at a "local optimum" and not reach a global optimum of $f(s)$.

allows occasional steps backward but requires overall forward movement.

C. Tentative Strategies

Backtracking



maintain list of states visited on current path only
[forgetting things that failed]
Problem: repeating past mistakes.

use a depth bound (n)

advantage - if a solution can be found in n moves or less,
backtrack will find it.

if not, will fail

iterative deepening:

→ try solving w/ depth n
if no solution, increase n

heuristics, informed search

- instead of trying moves in arbitrary or random order, try them according to "most promising" first, etc.

GraphSearch - Keeps track of a subtree of the solution space
each node has exactly one parent
(except root = initial state, with no parent)
as new nodes generated, either

- never seen before: use node that generated it as its parent
- seen before: choose between current parent + node that generated it.

[do this so that path from root to new node is min. cost or shortest

Note: this method can produce a
breadth-first search
or depth-first search
or best-first search

depending entirely on the order in which un-explored nodes are examined.

"Algorithm A" [which is really a class of evaluation functions]

Says :

nodes on the OPEN list should be stored in order from least to highest value of a function $f(n)$ having the form

$$f(n) = \text{Depth}(n) + \underline{h(n)}$$

→ estimate of distance from n to a goal node

Thus, $f(n)$ estimates the total distance from start to a goal node if going through node n.

Note: If $h(n)$ is chosen carefully enough to ensure that it never overestimates the distance from n to goal it is called an admissible heuristic (i.e., $h(n) \leq \underline{h^*(n)}$) and use of it with Algorithm A is known as Algorithm A* and guarantees a shortest path solution to a goal.

Note: If $h_1(n)$ and $h_2(n)$ are both admissible (i.e., $h_1(n) \leq h^*(n)$ and $h_2(n) \leq h^*(n)$ for all nodes n) actual distance (usually unknown)

Then $h_2(n)$ is more informed than $h_1(n)$ if $h_1(n) \leq h_2(n)$ for all n and provides better performance than $h_1(n)$ (i.e., fewer (or at least no more) node expansions to find solution)

Note $h(n) = 0$ in all cases is admissible
satisfies these conditions

\Rightarrow so $f(n) = \text{depth}(n) + 0$

provides breadth-first search

(because it ensures all nodes at level K will
be examined before any nodes at level $K+1$,
...)

Other stuff:

Agents : Performance Measures
Environment
Actuators
Sensors

deterministic,
stochastic

fully or partially observed...